




# Workshop: Forecasting time series

*VisCon, 12 October 2019*

digital natives

[unit8.co](https://unit8.co)



# About us

**Krzysiek & Kilian**



**Software Engineers @ Unit8**


Before we begin... Let's setup!

<https://github.com/unit8co/u8timeseries-viscon>


# Why Time Series Forecasting?

# Time Series Examples


Sales of new one-family houses, USA




US treasury bill contracts



Australian quarterly electricity production



Google daily changes in closing stock price



# The Big Picture

Everyone needs forecasting!



**Predicting  
electricity  
demand**


**Predicting  
product  
sales**

**Predicting  
taxi rides  
demand**

**... and  
many  
more**

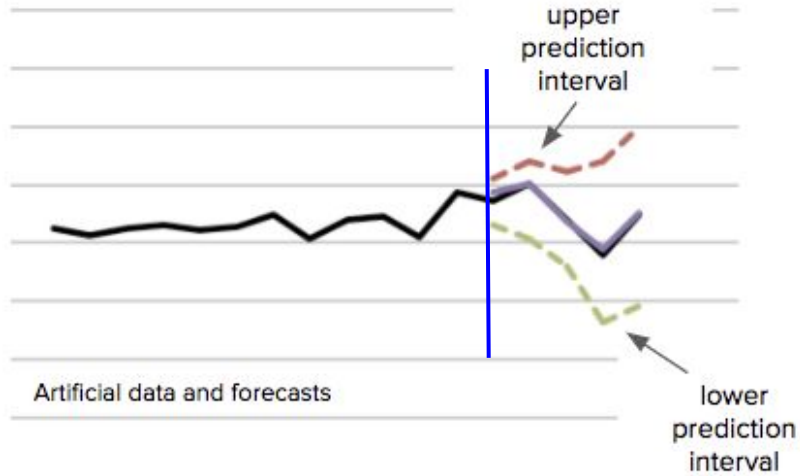
# Time Series Concepts

# Point Forecasts, Horizon, Residuals

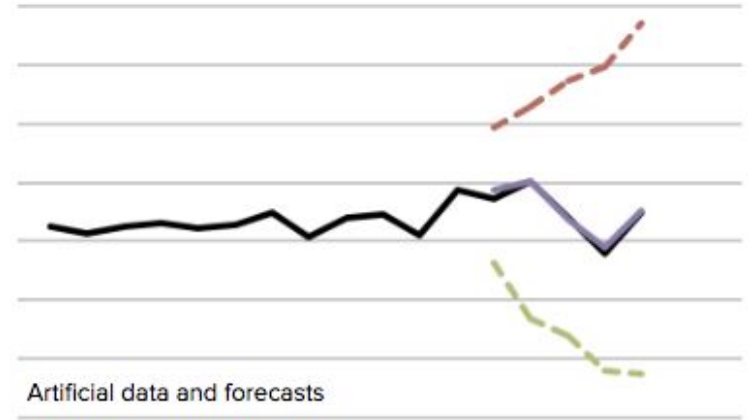




# Prediction Intervals




VS.

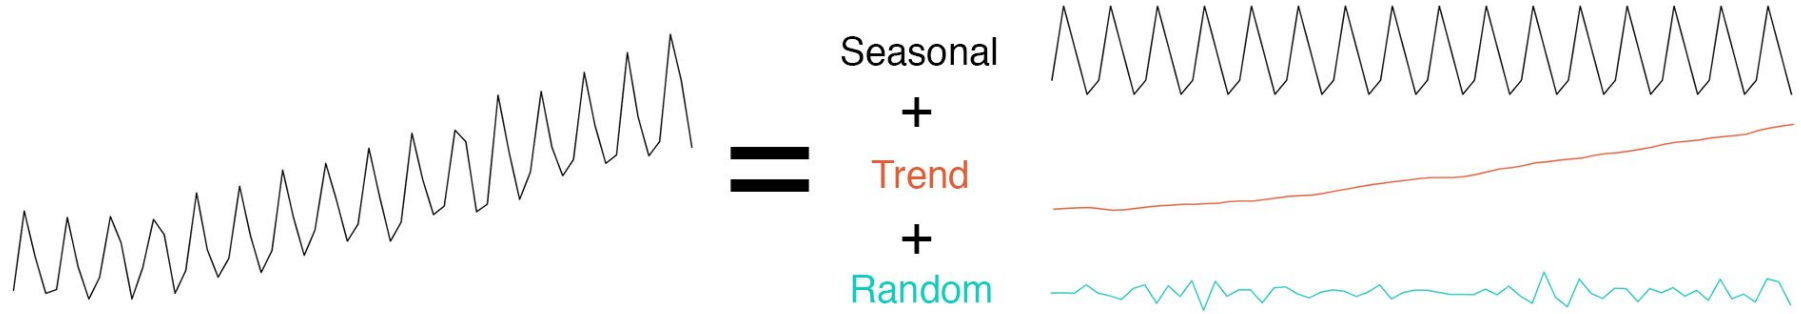


**We're 95% sure that the value is going to be within**

# Autocorrelation



# Seasonal Trend Decomposition




**Additive:**

Time series = Seasonal + Trend + Random

**Multiplicative:**

Time series = Trend \* Seasonal \* Random

# More complex trend & seasonality



# Time for Task 1


# The dataset

*RetailNZ.csv*

Monthly total clothing retail sales

Value in dollars

Timespan: May 1995 - September 2010



# Let's code!

Task 1 in *1\_explore\_and\_forecast\_task1-3.ipynb* notebook


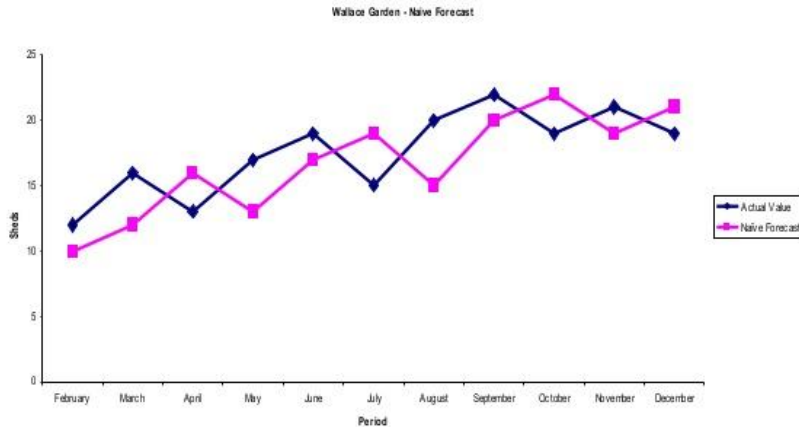
# Naive



# Naive

Just take the last value (or previous year value) as prediction of next value.

## Naïve Forecast Graph



# Exponential Smoothing


# Exponential Smoothing Intuition

Let's mix a weighted sum of:

- ★ **Previous value**
- ★ **Our previous estimation**

$$\hat{y}_{t+1|t} = \alpha y_t + (1 - \alpha) \hat{y}_{t|t-1},$$

Smoothing coefficients  
fitted to dataset



# Exponential Smoothing

There's lots of different methods in this family....

Holt's Method

Trend Component	Seasonal Component		
	N	A	M
	(None)	(Additive)	(Multiplicative)
N (None)	(N,N)	(N,A)	(N,M)
A (Additive)	(A,N)	(A,A)	(A,M)
A <sub>d</sub> (Additive damped)	(A <sub>d</sub> ,N)	(A <sub>d</sub> ,A)	(A <sub>d</sub> ,M)

Simple Exponential  
Smoothing (SES)

Holt-Winters'  
Method

# Exponential Smoothing

How to choose the model?

Follow ES modeling procedure

Better to automate the process!

```
library(forecast)
model <- ets(dataset)
future <- forecast(model, h=3)
```

`ets()` brute-forces all possible models to find the one that best fits the dataset

# ARIMA

# ARIMA = AR + I + MA

Let's mix together:

- ★ previous  $p$  values of the serie (AR)
- ★ previous  $q$  residuals of the serie (MA)
- ★ single or double differencing ( $d$ ) to get rid of seasonality & trend (I)  $y'_t = y_t - y_{t-1}$

Autoregressive AR( $p$ )      Moving Average MA( $q$ )

$$X_t = c + \varepsilon_t + \sum_{i=1}^p \varphi_i X_{t-i} + \sum_{i=1}^q \theta_i \varepsilon_{t-i}$$

Differencing I( $d$ )

Weights that need to be fitted on the dataset

# ARIMA = AR(p) + I(d) + MA(q)

How to choose p, d & q?

Follow ARIMA modeling procedure

Better to automate the process!

```
library(forecast)
model <- auto.arima(dataset)
future <- forecast(model, h=3)
```

`auto.arima()` brute-forces the search space of (p,d,q) that optimizes AIC, AICc or BIC value.



# Time for Task 2

# Error metrics

- Mean Absolute Percentage Error (MAPE)

$$M = \frac{100\%}{n} \sum_{t=1}^n \left| \frac{A_t - F_t}{A_t} \right|,$$

- Symmetric MAPE (sMAPE)

$$\text{SMAPE} = \frac{100\%}{n} \sum_{t=1}^n \frac{|F_t - A_t|}{|A_t| + |F_t|}$$

- Mean Absolute Scaled Error (MASE)

$$\text{MASE} = \text{mean} \left( \frac{|e_j|}{\frac{1}{T-1} \sum_{t=2}^T |Y_t - Y_{t-1}|} \right) = \frac{\frac{1}{J} \sum_j |e_j|}{\frac{1}{T-1} \sum_{t=2}^T |Y_t - Y_{t-1}|}$$

# Time for Task 3

**Which model is the best?**

# M3 Competition (2000) & NN3 (2006)

Can be expressed as Simple  
Exponential Smoothing (SES) with drift

- 3003 series
- Data from business, finance & economics
- Seasonality: yearly, monthly, daily, hourly
- Series length between 14 and 126

Method	MAPE	sMAPE	MASE
Theta	17.42	12.76	1.39
ForecastPro	18.00	13.06	1.47
B-J automatic	19.13	13.72	1.54
ETS	17.38	13.13	1.43
AutoARIMA	19.12	13.85	1.47

- Conclusions:
  - Complex methods not necessarily better than simpler ones
  - Methods ranking varies depending on accuracy metric & forecasting horizon
  - ML & Neural Networks unable to get comparable results

**u8timeseries**

# u8timeseries

- Specialized software written by Unit8
- Easy to use collection of forecasting models in Python
- Unified API, on top of *pandas*
- Auto-regressive and regressive models
- Utilities to manipulate time series

# Time for Task 4



# Ensembling

# M4 Competition (2018)

- 100K series
- Objective: compare ML vs statistical
- Point Forecast + Prediction Interval

Ranking	Type of Method	sMAPE	MASE	OWA
1st Ranked: Best	Hybrid	11.374	1.536	0.821
2nd Ranked: 2nd Best	Best Combining	11.720	1.551	0.838
8th Ranked	Best Statistical	11.986	1.601	0.861
19th Ranked	Comb Benchmark	12.555	1.663	0.898
25th Ranked	Best ML	12.894	1.682	0.915
37th Ranked	Naïve 2 Benchmark	13.564	1.912	1.000
48th Ranked	2nd Best ML	16.638	2.056	1.151

- Conclusions:
  - Combination of methods is the king - 12/17 of best methods
  - Surprising winner: hybrid approach ES-RNN submitted by Slawek Smyl from Uber
  - Pure ML methods performed poorly

# M4 Competition (2018)

Significant margin



Rank	Team	Affiliation	Method	sMAPE	MASE	OWA	Diff from Comb (%)
1	Smyl	Uber Technologies	Hybrid	11.37	1.54	<b>0.821</b>	-8.52
2	Montero-Manso et al.	University of A Coruña & Monash University	Comb (S & ML)	11.72	1.55	<b>0.838</b>	-6.65
3	Pawlikowski et al.	ProLogistica Soft	Comb (S)	11.84	1.55	<b>0.841</b>	-6.25
4	Jaganathan & Prakash	Individual	Comb (S & ML)	11.70	1.57	<b>0.842</b>	-6.17
5	Fiorucci, J. A. & Louzada	University of Brasilia & University of São Paulo	Comb (S)	11.84	1.55	<b>0.843</b>	-6.10
6	Petropoulos & Svetunkov	University of Bath & Lancaster University	Comb (S)	11.89	1.57	<b>0.848</b>	-5.55

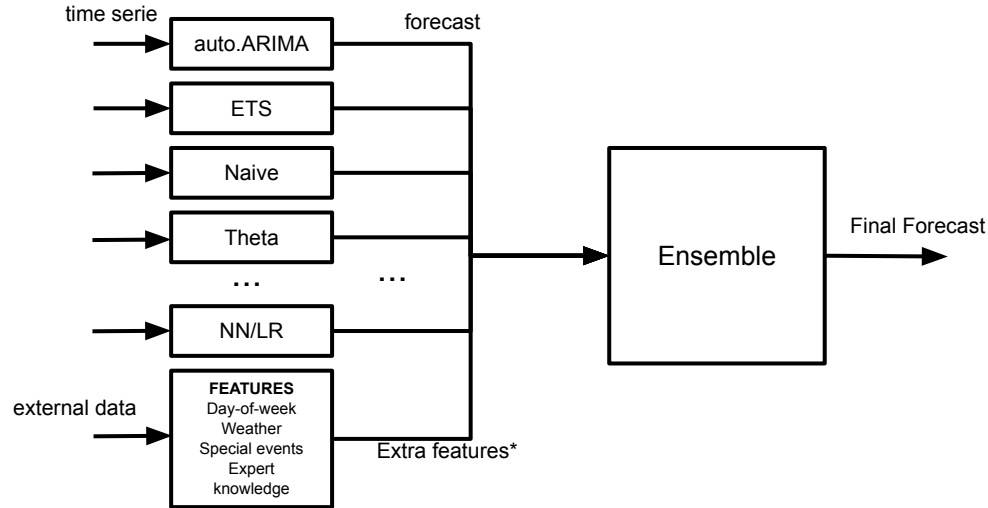
All submitted methods open sourced on GitHub:

<https://github.com/M4Competition/M4-methods>

All M4-Conference papers available online:

<https://www.mcompetitions.unic.ac.cy/presentations/>

# Ensembling



## Different ways to combine:

- Simple Average
- Weighted Average by accuracy
- Learned weights (LR, RF, XGBoost)

## Benefits:

- Improved accuracy due to additional information
- Safety in numbers - reduced risk of bad forecast
- Ability to incorporate external features


# Backtesting

# Backtesting

Concept: simulate predictions that would have been obtained historically with certain model(s)

1. Set prediction time  $t \leftarrow t_0$
2. Produce point prediction for a fixed time horizon
3. Move pointer  $t \leftarrow t+1$
4. Repeat

# Backtesting - test your model over time



aka Time Series Cross-Validation

# Time for Task 5



# Let's code!

*Workshop u8timeseries.* Jupyter notebook.

Tasks:

- Time series forecasting using u8timeseries
- Backtesting
- Ensemble predictions

# Forecasting Research

# In Search of Single Forecasting Model

## **Statistical methods work very well but:**

- Infeasible for millions of time series (Uber, Amazon)
- Require frequent retraining
- Do not benefit from cross-learning

## **Benefits of single ML model:**

- Difficult to train & get right, but easy to maintain
- Allow cross-learning & can accomodate for new time series

# Hybrid ES-RNN (Uber)

The idea is the following:

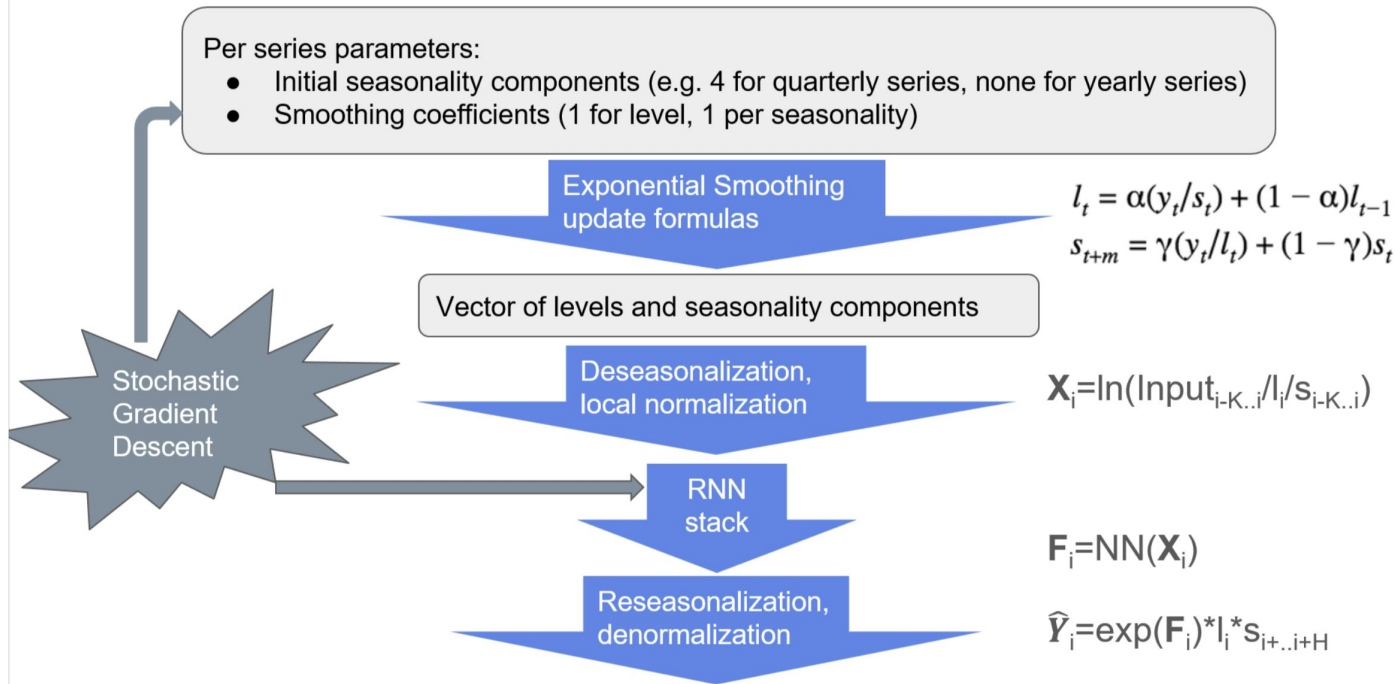
- Separate Exponential Smoothing (ES) model per serie
- One global LSTM network that learns on all of the series

## Ensemble of Specialists

- Train a pool of models on random subset of the series
- Final forecast for a particular serie is the average of top N models

# Hybrid ES-RNN (Uber)

## Dataflow



# DeepAR (Amazon)

- Available as a Service via AWS Forecast
- Also has some open source implementation <https://github.com/zhykoties/DeepAR>
- Used internally to forecast Amazon.com products demand & AWS utilization

## Highlights:

- One global LSTM model for all series
- Predicting likelihood instead of point forecast
- Covariates: item-dependent + time-dependent features

# Conclusions

- **Classical methods are hard to beat when you have:**
  - Sufficient history
  - Little external information
  - Few related time series
- **Start with classics and proceed with Combination of Methods:**
  - Reduced risk of bad forecast
  - Ability to incorporate external features
- **Experiment with emerging ML/DL methods:**
  - You know what you are doing
  - Lots of related time series
  - Cross-learning between series
  - One global model
- **Forecasting toolset:**
  - Much better in R (forecast, hts, tsfeatures, tsintermittent, etc..)
  - Bit rusty in Python (statmodels, prophet, etc)

# Bibliography

- M4-Conference Presentations  
<https://www.mcompetitions.unic.ac.cy/presentations/>
- Forecasting at Uber  
<https://eng.uber.com/forecasting-introduction/>
- Amazon DeepAR+  
<https://docs.aws.amazon.com/forecast/latest/dg/aws-forecast-recipe-deeparplus.html>
- Automatic Algorithms for Time Series Forecasting  
<https://www.slideshare.net/hyndman/automatic-time-series-forecasting>
- Exploring Time Series Feature Space  
<https://www.slideshare.net/hyndman/exploring-the-feature-space-of-large-collections-of-time-series>
- Facebook Prophet  
<https://research.fb.com/prophet-forecasting-at-scale/>
- Guru of Time Series Forecasting Homepage  
<https://robjhyndman.com/>
- DeepAR+ Open Source Implementation  
<https://github.com/zhykoties/DeepAR>
- ES-RNN Open Source Implementation  
<https://github.com/damitkwr/ESRNN-GPU>



Unit8™



Twitter: [unit8co](#)



Linkedin: [unit8.co](#)



Instagram: [unit8.co](#)

**follow us!**

unit8.co

**thank  
you**

